Losing Streaks in Ranked League of Legends: A Statistical Test of the Loser's Queue Hypothesis

NA Solo/Duo, Season 16: 144,623 matches, 1,032 players

Joseph X. Li

li.joseph@nifs.ac.jp

August 2026

Abstract

The "loser's queue" hypothesis holds that ranked matchmaking deliberately assigns worse teammates or better opponents to players on losing streaks, extending those streaks beyond what variance produces. We test this hypothesis against the complete season-16 ranked solo/duo history of 1,032 players sampled uniformly across the competitive ladder (144,623 unique matches, approximately 1.45M player–game records). Five falsifiable tests are applied, and all favor fair matchmaking. Observed longest losing streaks match player-matched Bernoulli null models with a mean z-score of -0.003; lag-1 win autocorrelation is -0.002; and teammate quality, once corrected for a leakage artifact that contaminates naive estimates, is flat in streak depth (.500–.503). To account for the small but statistically significant win-rate dip observed at deep streaks, we construct a causally disciplined model of player tilt from performance markers normalized against each player's own history, fitted on half the player population and evaluated exclusively on the other half. Decomposing the conditional win-rate curve into skill composition and this tilt state leaves residuals of +0.2 to +1.1 percentage points at every streak depth: observed win rates sit at or slightly above what fairness and player psychology jointly predict, which is the signature of Elo mean reversion and the opposite of what a punitive queue requires. Player-side effects are real but modest: fast re-queueing occurs in 19.4% of games following losses versus 15.9% following wins, and playing outside one's three most practiced champions costs 5.6 points of win probability, the largest controllable effect we measure. All methods are implemented in dependency-free Python, and all fitted models are evaluated on held-out players only.

Plain-language summary

A ranked game is engineered to be a coin flip; that is the purpose of matchmaking. Fair coins produce streaks that feel personal. Over a 250-game season, a 50%-win-rate player has a roughly 99% chance of at least one 5-loss streak (expected count 3.9) and an 11% chance of a 10-loss streak. The observed season-16 data reproduce these predictions almost exactly: 78.1% of players hit a 5-loss streak (77.0% predicted) and 10.8% hit a 10-loss streak (10.6% predicted). For calibration, the probability that a given game begins a 10-loss streak (about 1 in 2,048) exceeds the five-year fatal-accident rate of commercial aviation per flight (1 in 5.6 million; IATA 2021–2025) by a factor of roughly 2,700. The honest comparison class for a losing streak is stubbing a toe, not a plane crash. The widespread "my teammates get worse when I lose" experience is a measurement artifact: a duo partner shares your losses, so their naively measured win rate is mechanically depressed in exactly the games where you are streaking. Corrected, the curve is flat. What is real is player-side: players re-queue faster after losses, drift off their practiced champions while losing, and receive, from a fair Elo system, slightly easier games as their rating dips.

1 Introduction

Belief in loser's queue is widespread in the League of Legends community, sustained by anecdote, by informal analyses of personal match histories, and by the genuine psychological weight of long losing streaks. The claim is rarely stated precisely, which has allowed it to persist without direct confrontation with data. In its strong form, however, it is a mechanism claim, and mechanism claims make testable predictions.

Previous informal treatments have generally suffered from one or more of three defects: they evaluate streak probabilities against incorrect intuitive baselines rather than against null models matched to the player's own win rate and game count; they measure teammate quality with estimators contaminated by shared-game leakage; and they attribute performance variation to the matchmaker without accounting for player-side state. The present work addresses all three. We collect the complete season-16 ranked history of a ladder-stratified player sample, formulate fair matchmaking and loser's queue as competing hypotheses with distinct predictions, and apply five falsifiable tests. Because the data reveal a small win-rate deficit at deep streaks that a naive reading would attribute to the system, we additionally construct a causally disciplined model of player tilt and use it to decompose the conditional win-rate curve into composition, psychology, and residual. The residual is the quantity of interest: it is the room available for a system-side mechanism.

2 Definitions and predictions

Fair matchmaking. Conditional on the ratings it shows the player: (1) each game is assembled so that both teams have approximately 50% win probability, and (2) game outcomes are independent given ratings; the system maintains no state that tracks streaks and uses none in team assembly.

Loser's queue (strong form). The system deliberately assigns worse teammates or better opponents following losses, extending losing streaks beyond variance.

Loser's queue (weak form). "My streaks feel engineered." This is a claim about experience rather than mechanism, and it can be true while the strong form is false; part of the present work is explaining why fair systems reliably produce that feeling.

Metric	Fair predicts	LQ predicts	Observed (S16)
Longest streak vs matched null	$z \approx 0$	$z \gg 0$	$z = -0.003$
$P(\text{win} \mid \geq k \text{ losses})$	$\approx$ baseline$^+$	below baseline	-1.7 pts at $k \geq 5^*$
Lag-1 autocorrelation	$\approx 0^-$	clearly positive	-0.002
Teammate WR vs streak	flat	declines	flat (.500–.503)
Residual after tilt account	$\gtrsim 0$	large negative	+0.2 to +1.1 pts

Table 1: Predictions versus observations. $^+$Elo mean reversion pushes the conditional win rate slightly above baseline. *Fully absorbed by the composition and tilt accounts (Figure 2).

3 Data and methods

3.1 Collection

We sample NA ranked solo/duo (queue 420) stratified across the ladder: 36 players from each of the 28 tier–divisions from Iron through Diamond, plus 8 each from Master, Grandmaster, and Challenger, for 1,032 seeded players. For each player we collect the complete season-16 ranked match history (season reset 2026-01-08), paginated through the match-v5 endpoint with the `startTime` filter. The resulting dataset contains 144,623 unique matches, each storing all 10 participants. Collection is checkpointed and resumable, respects the API rate limits with a client-side sliding-window limiter, and retains the full raw match records. Descriptive tests use all qualifying players; every fitted model uses a 516/516 player split, with all reported metrics computed exclusively on players the fit never saw.

3.2 Null models and descriptive tests

The headline test compares each player's observed longest losing streak against a matched null: for each player with at least 20 games ($n = 785$), we simulate 2,000 seasons of n_i Bernoulli(w_i) games at the player's own win rate w_i and game count n_i, and compute the z-score of the observed longest streak against the simulated distribution. This construction controls for the two features that mislead intuition, namely sub-50% win rates and large game counts. Analytic streak-occurrence probabilities provide an independent cross-check of the simulation. Supporting tests are the pooled lag-1 autocorrelation of the win indicator, conditional win rates by streak depth with Wilson 95% intervals, session structure (sessions split on 45-minute gaps), and teammate quality by streak depth. We note that fair Elo-style matchmaking is not an i.i.d. coin: visible rating mean-reverts around true skill, which makes outcomes slightly anti-correlated and predicts marginally *fewer* long streaks than the coin, together with a conditional win rate that *rises* with streak depth as the system serves easier games to a dipping rating.

3.3 The tilt model

Per game we compute three performance markers: deaths per minute, CS per minute, and KDA. Each marker is z-scored against the player's own baseline for the same outcome, so that losses are compared only with that player's other losses; in-game statistics are contaminated by the game's result, and outcome conditioning removes this channel. Baseline cells require at least 5 games, with fallback to outcome-only cells and a variance floor. The per-game sloppiness index is $s = \text{mean}(z_{\text{dpm}}, -z_{\text{cspm}}, -z_{\text{kda}})$.

All model features are computed on an expanding window: the features for game i use games $1..i-1$ only. This discipline is quantitatively important. With full-season baselines, which admit look-ahead, the apparent tilt gradient across state deciles is 3.9 points of win probability; with expanding windows it is approximately 1.2 points, and approximately 0.8 after champion-familiarity normalization. Informal analyses that normalize against full-history averages contain this artifact.

The pregame state entering game i comprises an exponentially weighted moving average of prior sloppiness ($\alpha = 0.3$), the loss streak entering the game (capped at 5), and a deep-session indicator (fifth game or later) interacted with the previous outcome. Player skill enters as the logit of a shrunken expanding win rate, $\tilde{w}_i = (W_i + 10)/(N_i + 20)$, an empirical-Bayes estimate with 20 pseudo-games at the population mean.

3.4 Weight determination

The weights were fitted, not chosen; our own starting intuitions were demoted by the data. The model is a ridge-penalized logistic regression,

$$P(\text{win}) \;=\; \sigma\Big(\beta_0 + \textstyle\sum_j \beta_j x_j\Big), \qquad \hat\beta \;=\; \arg\max_\beta \;\ell(\beta) - \tfrac{\lambda}{2}\|\beta\|^2, \quad \lambda = 10^{-3},$$

with continuous features standardized on the training half. The maximum-likelihood estimate is obtained by Newton–Raphson iteration,

$$\beta \leftarrow \beta + \big(X^\top W X + \lambda I\big)^{-1}\big(X^\top(y-p) - \lambda\beta\big), \qquad W = \operatorname{diag}\big(p_i(1-p_i)\big),$$

with damped steps, iterated to convergence. Fitting uses 516 players (69,789 games); all reported metrics come from the disjoint 516 (59,967 games). Feature selection was performed by ablation on held-out log-loss, evaluating each candidate feature both solo and by drop-one. The final standardized weights are given in Table 2.

Feature	$\hat\beta$	Note
player skill (logit shrunk WR)	+0.065	who you are
sloppiness EWMA	−0.023	the tilt core; strengthened when junk was pruned
streak entering	−0.007	loss-recency block
deep session × prev loss	−0.006	

Table 2: Final model. Features killed by ablation: fast re-queue following a loss (its coefficient is repeatedly *positive*, confounded with duo play and confidence), off-schedule hours, and all session-context features (session net score, last-game sloppiness, closeness of the previous loss, duo continuation), each of which overfit; the EWMA already summarizes session context. Our starting intuition, that previous outcome and streak should be weighted most heavily, did not survive: performance history carries more held-out signal than outcome history.

The derived tilt coefficient $T = P(\text{win} \mid \text{state})/P(\text{win} \mid \text{best-case state})$ has decile-binned expected calibration error below 1.5% in the mid-range and a Platt recalibration slope of 1.44 on session-final games: the model *underestimates* tail spread by roughly 40%, a conservatism we retain deliberately. Tilt-flagged session enders realize .478 against .491 predicted.

4 Results

4.1 Test 1: longest streak versus matched null

The mean z-score of observed longest losing streaks against the matched nulls is -0.003 (median -0.157). Streak incidence agrees with the null throughout: 5-loss streaks are observed in 78.1% of players against 77.0% expected; 8-loss streaks in 25.7% against 29.2% (fewer than fair); 10-loss streaks in 10.8% against 10.6%. The slight deficit at depth 8 is consistent with the mean-reversion signature discussed above: fair matchmaking produces marginally fewer long streaks than an i.i.d. coin.

4.2 Test 2: conditional win rate and the tilt-accounted decomposition

The observed conditional win rate declines gently with streak depth, from .4999 unconditioned to .4833 $[.4690, .4977]$ at $k \geq 5$ ($n = 4{,}645$), as shown in Figure 1. At this sample size the dip is statistically real. It is also small, and it is directionally *opposite* to the Elo-drift fair null, which rises to approximately .59 at $k \geq 5$.

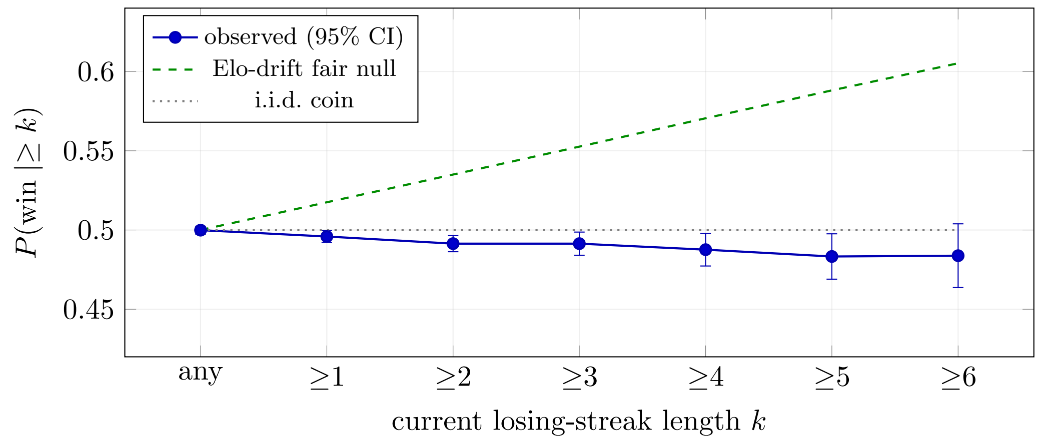

Figure 1: Conditional win rate by streak depth, all 1,032 players (145,816 conditioning games).

To determine whether this dip leaves room for a system-side mechanism, we fit two logistic accounts on the training half of players, neither of which sees streak information: a *composition* account (player skill only; streaks over-sample weaker players) and a *composition + tilt* account (adding the causal tilt state of Section 3.3). Their predictions are compared with observed win rates at each streak depth on held-out players only (59,967 games) in Figure 2. The residual at every depth through $k \geq 5$ is positive, +0.2 to +1.1 points ($k \geq 6$: -0.5 points, $n = 950$, with a $\pm$3-point interval). A loser's queue requires large negative residuals here. The small positive residuals are instead consistent with the uncompensated portion of the mean-reversion tailwind. After accounting for who is on streaks and the psychology they bring, the room remaining for a punitive queue is approximately zero.

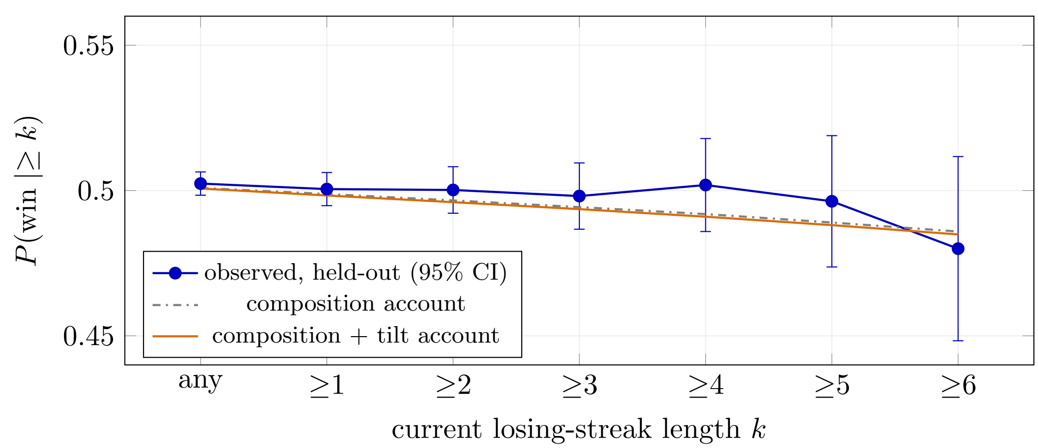

Figure 2: Observed conditional win rate against the composition and composition + tilt accounts, held-out players only. Observed values sit at or above the full account at every depth through $k \geq 5$.

4.3 Test 3: autocorrelation

The pooled lag-1 autocorrelation of the win indicator is -0.002. Fair matchmaking predicts a value near zero with a slight negative bias from mean reversion; any streak-extending mechanism requires it to be clearly positive. It is not.

4.4 Test 4: sessions and queueing behavior

Sessions were split on gaps of 45 minutes or more (median length 2 games, $p_{90} = 4$). Fast re-queueing (within 5 minutes) occurs in 19.4% of games following a loss against 15.9% following a win. The stopping decision itself is a recency reflex rather than ledger management: in aggregate, $P(\text{stop} \mid \text{loss}) = .387$ is indistinguishable from $P(\text{stop} \mid \text{win}) = .386$, and session net score does not modulate stopping. The structure lies elsewhere. After the first game of a session, players stop *more* often following a win; the stopping hazard peaks at the second consecutive loss (.408); and at three or more consecutive losses it falls below the after-win baseline (.369), consistent with self-selection of committed players into continued play.

4.5 Test 5: teammate quality and the leakage artifact

For every game we compute the mean in-sample win rate of the player's four assigned teammates, bucketed by the player's current streak. The naive estimator, which excludes only the current match from each teammate's record, declines monotonically: .5064, .5031, .4998, .4958, .4937 for streaks 0/1/2/3/4+. This is a textbook loser's-queue signature, and it is an artifact. Duo partners appear in many of the focal player's games, and their measured win rates include the games they lost together, so conditioning on the focal player's streak conditions on the teammate's own recorded losses. Excluding every shared match from both numerator and denominator yields .5029, .5032, .5013, .4996, .5016 (up to $n = 46{,}428$ per bucket): flat. Informal match-history analyses replicate the naive estimator and inherit its bias in full.

5 Discussion

Three observations organize the results. First, every direct test of the strong hypothesis returns the fair-matchmaking prediction, and the one statistically real departure from the i.i.d. baseline, the win-rate dip at deep streaks, carries the wrong sign for a punitive queue relative to the Elo null and is fully absorbed by composition and measured tilt. Second, the player-side effects that do exist are individually small. The pregame-predictable component of tilt is worth roughly one to two points of win probability, whereas champion familiarity is worth 5.6 points, several times larger than the tilt effect it is often mistaken for; part of the familiarity exposure is itself downstream of tilt, since players on streaks drift off their practiced champions three to four points more often. Third, the weak form of the hypothesis is explained rather than dismissed: fast re-queueing after losses clusters defeats into heated sessions, and within-period performance excursions are considerably larger than their pregame-predictable component, so the experience of engineered streaks is real as experience while the mechanism claim fails.

We note one further interpretive point. Because an Elo-anchored system demotes a player until their average performance, tilt episodes included, wins 50% of games at that level, a player's rank already capitalizes their characteristic tilt. The marginal within-rank effects measured here are therefore small by construction, and they do not bound the counterfactual effect of tilt on the rank a player equilibrates to. That counterfactual is not identifiable from single-season cross-sectional data; we outline the required multi-season trajectory design in the repository documentation.

Limitations. The sample covers one region and one season; 1,032 players is large for the questions asked but is not a census, and apex tiers are thin ($n = 24$). Seeding from the mid-August ladder biases first-half-season form slightly upward (a regression artifact measured at approximately 0.7 points). Teammate win rates are in-sample. The tilt model is predictive rather than interventional. The headline results are robust to the session-gap (30/45/60 minutes) and streak-threshold choices.

6 Conclusion

We tested the loser's queue hypothesis against the complete season-16 ranked history of 1,032 ladder-stratified players. Observed streak statistics match player-matched Bernoulli nulls; autocorrelation is consistent with zero; corrected teammate quality is flat in streak depth; and the conditional win-rate curve, decomposed into skill composition and a causally measured tilt state, leaves no negative residual for a system-side mechanism to occupy. Overall, the present results support the conclusion that losing streaks in ranked solo/duo are ordinary variance shaped by fair matchmaking, that the streak-associated performance deficit is player-side and modest, and that champion familiarity, not tilt and not the matchmaker, is the largest factor under a player's direct control. Future work will be directed toward cross-region replication (EUW alongside NA), behavioral intensity markers (ping and ability-cast telemetry) as a sharper sensor for tilt than outcome statistics, and, given census-scale access, the lobby-level tilt differential and the multi-season rank-trajectory designs that the present dataset cannot support.

Statistical methods

Wilson score intervals for all proportions; Monte Carlo matched-null simulation (2,000 seasons per player) with analytic cross-checks; pooled lag-1 autocorrelation; within-player, within-role, outcome-conditioned z-scores with thin-cell fallback and variance floors; exponentially weighted moving averages; expanding-window (causal) feature construction; empirical-Bayes shrinkage via pseudo-counts; ridge-penalized logistic regression by Newton–Raphson on standardized designs; player-level train/test splits; solo and drop-one ablation on held-out log-loss; permutation tests (2,000 shuffles); expected calibration error and Platt-slope recalibration; inverse-Simpson effective counts for champion pools; right-censored hazard analysis for session stopping; and duo detection by cross-game teammate intersection. All methods are implemented from scratch in dependency-free Python 3.9+.

Data availability

The collection and analysis code (`collect.py`, `null_model.py`, `analyze.py`, `tilt.py`, `tilt_coeff.py`, `refit.py`, `streak_account.py`, `stopping.py`, `profiles.py`, `tier_tilt.py`), the aggregate results files, and this report are available in the public repository (link to be added). Match data were obtained from the Riot Games API (league-v4, match-v5) in August 2026. This report is not endorsed by Riot Games.

Sources

IATA Annual Safety Report 2025 (five-year average fatal-accident rate, 2021–2025). BlenderTimer toe-stubbing survey ($n = 2{,}866$; self-reported, treated as directional). Cleveland Clinic, “Stubbed Toe.”